

HiringCompass: LG6 · Assessment graph (F6)

One run per application and compatible versions · thread_id = application + versions · chained after a valid F5 result · never reruns upstream reasoning

Legend

LLM node (retry policy)

Deterministic / tool node

Conditional edge (router)

HITL interrupt (pause / resume)

END state

Persistence

LangGraph run — assessment graph

```
graph TD; Worker[Worker  
claims work item after F5] --> START([START]); START --> load_upstream[load_upstream  
F1 · F2 · F5 results]; load_upstream --> router1{same job and  
rubric version?}; router1 -- no --> END1([END  
reevaluation needed]); router1 -- yes --> merge_evidence[merge_evidence  
CV + interview evidence]; merge_evidence --> generate_assessment[generate_assessment  
strengths · gaps · follow-ups]; generate_assessment --> router2{conclusions  
traceable?}; router2 -- "no · regenerate" --> generate_assessment; router2 -- exhausted --> END2([END  
failed (retryable)]); router2 -- yes --> persist_assessment[persist_assessment  
assessment · evidence links]; persist_assessment -- write --> PG2[(PostgreSQL  
assessment · reports · decision)]; persist_assessment --> render_report[render_report  
printable HTML]; render_report --> HITL{{HITL: final decision}}; HITL --> record_decision[record_decision  
actor · reason · time]; record_decision --> END3([END  
decision recorded]); PG1[(PostgreSQL checkpoint  
checkpoint after every node)] -.-> PG2;
```

The flowchart illustrates the assessment graph process. It begins with a **Worker** claiming a work item after F5, leading to a **START** node. The process then moves to **load_upstream** (F1 · F2 · F5 results). A decision point checks if the job and rubric version are the same. If not, it leads to **END reevaluation needed**. If yes, it proceeds to **merge_evidence** (CV + interview evidence). This is followed by **generate_assessment** (strengths · gaps · follow-ups). Another decision point checks if conclusions are traceable. If not, it loops back to regenerate. If exhausted, it leads to **END failed (retryable)**. If yes, it proceeds to **persist_assessment** (assessment · evidence links). This node has a **write** connection to a **PostgreSQL** database (assessment · reports · decision) and leads to **render_report** (printable HTML). The process then reaches a **HITL: final decision** node, followed by **record_decision** (actor · reason · time), and finally **END decision recorded**. A **PostgreSQL checkpoint** (checkpoint after every node) is also shown, connected to the main PostgreSQL database.